MEMBRIA.AI · PARTNERSHIPS

The experience layer for *AI reasoning*.

Membria records what agents decide and what happened next, so every agent improves without changing model weights or grading itself — an auditable alternative to recursive self-improvement. We partner with technology and infrastructure providers to bring it to more teams, clouds and data centres.



Michael Aprossine

CEO, Co-Founder — Membria.AI by Actiquest Labs

PARTNER DECK

2026 · v1.0

COMPANY

Who We Are



Actiquest Labs

Applied AI Systems Company

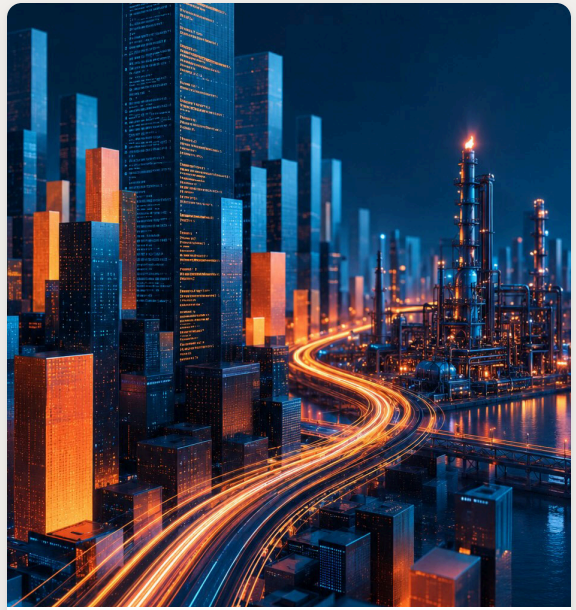
Practical automation, intelligent workflows and enterprise-grade AI — built for real results.



Membria.AI

AI Operating Environment

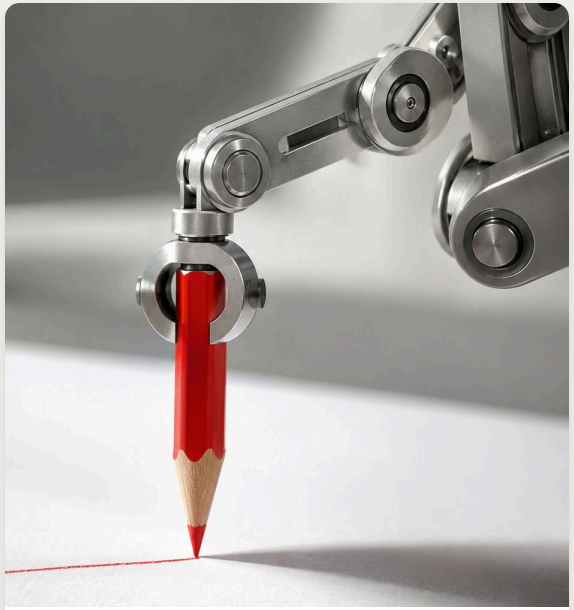
Our core product — agents, memory, knowledge and verification combined into one enterprise-ready AI platform.



Vertical Expertise

Code · Fintech · Government · Audit · EPC

We specialize in complex domains where AI needs precision, context and process understanding.



How We Work

Applied AI Engineering

We ship enterprise AI projects that stress-test and fund Membria's core — concept to production.

COMPANY PROFILE

Actiquest Labs at a Glance

Key facts, production footprint and enterprise AI capabilities.

10+ years enterprise product, AI engineering and venture-building

4 verticals — Fintech · Government · Audit · Engineering/EPC

100+ seats deployed across production AI use cases

METRIC	VALUE
Founded	February 2023 · California, USA
European Presence	Belgium entity · relocation to Luxembourg planned
Core Product	Membria.AI — memory & experience layer for AI agents
Production Verticals	SDLC · Fintech/Banking · AI Accounting/Audit · Engineering/EPC
Deployment	Cloud SaaS · self-hosted · on-premise enterprise
Clients & Partners	OES · Wowcube · Hiscox · K-Telco · Kazakhstan Government · industrial enterprises across EMEA/CIS
Technology Partners	NVIDIA · Amazon Web Services · LuxProvide
Team	core team of 4 · 20+ years enterprise software & audit · 9 years applied AI — finance, government, industrial, deep-tech

OPPORTUNITY

Why partner *with Membria*

1. **Agents became the workforce**
   Claude Code, Codex and Cursor do real daily work in code, engineering and finance — but they ship without memory. Membria adds it.

2. **One integration, every agent**
   MCP plugs a memory plane into any agent on the market — technology partners connect once instead of building per-vendor integrations.

3. **Sovereign & regulated by design**
   EU AI Act and ISO 19650 require defensible decisions and audit trails. Membria runs cloud, self-hosted or on-premise — a fit for sovereign and regulated infrastructure.

4. **Experience compounds with every partner**
   Nobody owns the decision → outcome graph yet. Every deployment strengthens the shared knowledge graph that all partners build on.

PRODUCT

Membria *Reasoning Graph* Workforce

INSTITUTIONAL KNOWLEDGE



Memory

- ✓ **Decisions** — every choice captured with its reasoning
- ✓ **Consistency** — behavior chains keep teams aligned
- ✓ **Skills** — wins distilled into reusable know-how
- ✓ **NegativeKnowledge** — failures become rules

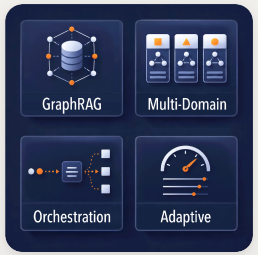
MULTI-AGENT HARNESS



Reasoning

- ✓ **Mesh** — many agents work in parallel
- ✓ **Coordination** — agents hand off via A2A
- ✓ **Routing** — each task to the right model
- ✓ **Quality gates** — verification blocks bad output

DOMAIN-AGNOSTIC FOUNDATION



Graph

- ✓ **GraphRAG** — semantic search over a property graph
- ✓ **Multi-domain** — one schema across every domain
- ✓ **Orchestration** — work runs as a dependency graph
- ✓ **Adaptive** — memory that strengthens with use

HOW IT WORKS

One loop, *every stage*

Not a memory bolted onto one step — every shipped outcome feeds the next decision.

01

Plan

Recalls decisions and outcomes from similar work — the agent starts from what already worked.

02

Execute

Injects validated skills and blocks known anti-patterns as the work is being done.

03

Review

Checks the result against past failures and standards — deterministic critics, auditable output.

04

Ship

Records the real outcome and links it back to the decision that caused it.

↻ **Every outcome feeds the next decision.** One MCP connection: `npx @membria/cli init`

WHY MEMBRIA?

Core Capabilities

The full stack — from memory and reasoning to deployment and governance



AI Agent Orchestration

Multi-agent BBS mesh, WorkGraph DAG, A2A protocol coordination



Causal Memory

Decision / Outcome / NegativeKnowledge graph with time-decay and compounding



Verification Engine

Cross-document consistency checks, rule-based normative validation, quality gates



Graph Intelligence

Memgraph property graph, symbol-level CodeGraph, policy graph with Hebbian co-activation



LLM Integration

Claude / GPT factory, structured output, behavior chains, context injection



Product Engineering

Full-stack: FastAPI + React + Node.js, MCP protocol, Telegram, Web Dashboard



Domain AI Systems

Vertical-specific schemas, critics, standards compliance (ISA / API / ASME / IEC)



Security & Governance

PolicyEngine, RBAC, OIDC/SSO, prompt injection defense, secrets management

Bottom line: Generic AI forgets. Membria remembers, reasons, and governs.

ARCHITECTURE

Architecture *& integration*

One MCP endpoint in front, one graph core behind, any model and any infrastructure underneath.

Agents & clients	Claude Code · Codex · Cursor · Telegram · Web dashboard · CLI
Integration	MCP — one endpoint · A2A agent coordination · Memory Contract API · Python client library · Kafka event bus
Membria core	Causal memory: Decision · Outcome · NegativeKnowledge · Skill · Multi-agent mesh & quality gates · Verification engine · WorkGraph · CodeGraph · GraphRAG
Data & models	Memgraph property graph · Qdrant vector search · Postgres · S3-compatible object storage · LLM gateway: Claude · GPT · GigaChat · on-prem open-weight (vLLM)
Deployment & governance	Cloud SaaS · Self-hosted · On-premise · Offline signed bundles for closed networks · OIDC / SSO · RBAC · OPA policy · audit log

Connect in one command: `npx @membria/cli init` — customer data stays in a private per-tenant graph; nothing is published unless you choose to.

FOCUS

One core, *many domains*

Verticals are distribution of one product — not divergence. The same graph schema ships verbatim in every codebase.

PRODUCT	WHO BUYS	DELIVERY	STATUS
Membria CE	solo developers on Claude Code / Codex / Cursor	cloud SaaS · zero-install MCP	SAAS · CODE.MEMBRIA.AI
Membria Enterprise	engineering & product teams	self-hosted platform · own UI	DEMO · DESIGN PARTNERS
AI-EPC	EPC contractors & design institutes	vertical app on the same core	LIVE DEMO · EPC.MEMBRIA.AI
Glass	accounting & audit teams	vertical app on the same core	DESIGN-PARTNER STAGE

Proof of substrate: Decision · Outcome · NegativeKnowledge · Skill — one schema, one causal experience loop, shared verbatim across every codebase. A vertical competitor can't follow us across domains; we can.

PROJECTS

What we are *building*

Six projects on one memory core — from developer tools to regulated finance and public safety.

PROJECT	CUSTOMER / DOMAIN	WHAT WE BUILT	STATUS
Membria CE	solo developers · Claude Code, Codex, Cursor (API to B2B)	remembers every decision and what happened next, so coding agents stop repeating mistakes	SAAS
Membria CLI	engineering teams · inside the customer perimeter	one delivery pipeline for a team's agents, from spec to release, with gates, policy and audit	ON-PREMISE
AI-EPC	EPC contractors & design institutes · Worldwide	turns a tender folder into a checked, issued engineering package across 17 disciplines	SAAS
Crypto broker platform	licensed virtual-asset broker · Central Asia	data bus and decision memory that connect execution and compliance	ON-PREMISE
Bank automation	developer-owned bank · CIS	regulatory report specs, mortgage checks and escrow rules, drafted and validated by agents	ON-PREMISE
C2 Center AI	government emergency response	early warning and road-closure control, with every decision approved by a person	ON-PREMISE

Two delivery models: SaaS to start in minutes, On-Premise inside the customer's perimeter. Client names are withheld.

CASE STUDY · 01

Membria CE *learns from reverts*

Solo developers · Claude Code, Codex, Cursor (API to B2B) · SaaS

1. The situation

A developer asks the agent to add caching to an API client. The agent reaches for an in-memory LRU — the same choice that broke under multi-worker deploys and was reverted two weeks ago.

2. How it works

Before choosing, the agent asks Membria for decision context. It gets the reverted decision with its lesson (per-process cache, inconsistent across workers) and a signal from the commons: most projects that adopt this pattern later remove it.

3. The result

The agent picks a shared cache and records the decision with its alternatives. When the change survives, the outcome attaches itself from git and CI — nobody writes a post-mortem. The agent gets better without touching model weights and without grading itself: an alternative to recursive self-improvement, grounded in what really happened.

96.5K

production repos mined

~570K

Skills + AntiPatterns

1

command to connect

Stack & partners: MCP · Claude Code · Codex · Cursor · code.membria.ai

CASE STUDY · 02

Membria CLI *agents as one team*

Engineering teams · inside the customer perimeter · On-Premise

1. The situation

A team of six has to add rate limiting to the public API. Instead of six separate agent sessions, the task enters one pipeline: Researcher, Planner, Developer, Reviewer and Gate.

2. How it works

The Researcher surfaces a per-process counter the team shipped and reverted last year. The Planner builds the plan against CodeGraph and adds a load test. An engineer's own Claude Code takes the Developer seat over MCP.

3. The result

The Gate blocks the first attempt for missing load-test evidence and passes the second. Every step has an actor, a model, a policy decision and an audit record, and the team graph never leaves the perimeter.

8

pipeline phases, spec to release

5

agent roles, one blocking gate

100%

of the team graph stays in the perimeter

Stack & partners: MCP · OIDC / SSO · OPA · any LLM provider

CASE STUDY · 03

AI-EPC *from tender to issue*

EPC contractors & design institutes · Worldwide · SaaS

1. The situation

An EPC contractor uploads a tender folder: specifications, drawings, vendor catalogues, contract terms. This is where requirements get missed, and a late clause turns into rework — 28% of budget overruns.

2. How it works

AI-EPC reads every file, scanned pages included, into a requirements register where each line cites its source clause. It builds the design level by level: models propose, rules, codes and catalogues accept or reject, and unknown values become questions to the customer.

3. The result

Each discipline lead approves documents at a gate, and a change in one place reaches every dependent sheet. The contractor issues a consistent package with an evidence trail behind each document. Engineers sign; the system never stamps.

17

engineering disciplines

20,000+

standards in the graph

5

approval gates, G0 to G4

Stack & partners: DXF · DEXPI · IFC · ISO / IEC / ASME / API / GOST · epc.membria.ai

CASE STUDY · 04

Crypto broker *one data bus*

Licensed virtual-asset broker · Central Asia · client anonymised · On-Premise

1. The situation

A licensed broker runs execution and compliance as two parallel tracks built by two vendors. An operation must not execute until compliance has decided on it, and personal data must stay out of shared systems.

2. How it works

Membria owns the data bus and decision memory between the tracks: a Kafka backbone with 38 schema-validated event types and a graph of 42 object types. A client's status and every operation decision travel to the other track over the bus.

3. The result

A compliance officer asks in plain language why an operation was blocked and gets the full chain — screening result, limit breaches, approval steps — with the source events. Both vendors integrate against one documented contract.

38

event types on one bus

376

automated tests passing

~6 wks

from design to first full cycle

Stack & partners: Kafka · Memgraph · Qdrant · FastAPI · Python

CASE STUDY · 05

Bank *reporting & mortgage*

Developer-owned bank · CIS · regulatory reporting, mortgage and escrow · On-Premise

1. **The situation**

Methodologists spend 2–5 business days on each central-bank report specification. Mortgage review is a 40+ field manual checklist, and escrow specs are copied from the previous project.

2. **How it works**

A reporting assistant writes the specification from a free-form task, using graphs of the bank's object model and the regulator's forms. Skills keep the mappings that worked; rejected mappings are never proposed again. A mortgage assistant checks each application against project and escrow data.

3. **The result**

A specification is ready in 1–4 hours instead of 2–5 days, and when the regulator changes a form, every affected specification is listed at once. The bank picks its own model — Claude, GigaChat or on-prem Qwen — and data stays inside its security boundary.

1–4 h

per specification, down from 2–5 days

3

rollout phases, about 20 weeks

4

domain graphs: model, regulatory, compliance, escrow

Stack & partners: Graph memory · Claude / GigaChat / on-prem Qwen · web, Telegram, CLI

CASE STUDY · 06

C2 Center *decisions in minutes*

Government emergency response · On-Premise

1. **The alert**
   Drone LiDAR shows slope movement near a bridge well above the seasonal limit, while river levels rise. The duty officer needs one decision-ready picture, not five separate systems.

2. **How it works**
   The Duty agent matches the finding with the object's history in the knowledge graph and raises a risk alert with a confidence score and cited sources. The Planner offers three response options, Routes builds detours, and quality control blocks any option that repeats a recorded unsafe decision.

3. **The result**
   The chief of staff approves, and only then do alerts go out in two languages by SMS, Cell Broadcast and sirens. On closed roads, edge-AI cameras send plate events instead of video, with no vehicle-owner registry; the full trail stays on-premise.

5
AI agents, humans decide

100%
on-premise processing

0
personal-data registries

Stack & partners: LiDAR · IoT sensors · MQTT with mutual TLS · edge-AI cameras

TEAM

The people *behind it*

Founders and operators with two decades across enterprise IT, IoT, fintech and applied AI.



Michael Aprossine

CEO · CO-FOUNDER

Serial founder with 20+ years of product leadership across enterprise IT, IoT, mobility and blockchain. Started in IT consulting at a Microsoft Gold Partner, co-founded Apla/eGaaS and sold the mobility startup DrivePoint to Hyundai. Owns Membria's product vision and system architecture.

Focus: product vision, system architecture, AI platform
Previously: Apla/eGaaS · DrivePoint (now Hyundai Mobility) · Slotlogic IoT/SaaS · Microsoft Gold Partner

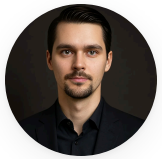


Philippe Khomenok

COO · CO-FOUNDER

Commercial builder who turns new markets into revenue. Scaled online sales of cable and rolled-metal products to $10M a year from scratch, co-founded the Grown Capital fund ($5M raised) and, as CBDO at Advcash, launched new products and entered new markets. Leads Membria's go-to-market and partnerships.

Focus: go-to-market, partnerships, growth
Previously: Advcash (CBDO) · Grown Capital (co-founder)



Mike Keer

CTO

AI architect with deep expertise in language models and computer vision. Leads Membria's engineering and designs its memory and reasoning systems. Has worked on joint projects with Luxinnovation, the European Innovation Council and others.

Focus: memory & reasoning systems, LLMs, computer vision
Previously: joint projects with Luxinnovation and the EIC



Mayya Aprosina

UI/UX

Product designer with 6+ years leading cross-functional teams through the full product cycle — research, journey mapping, design systems and front-end build in React and Next.js. Founder of the SiteSynth web studio and the designer behind ScoreSynth and Membria.

Focus: UI/UX, design systems, front-end
Previously: SiteSynth (founder) · Idalko (digital specialist partner)

INFRASTRUCTURE PARTNERS

What we need *from infrastructure*

Agent memory is RAM-bound and latency-critical — the profile a sovereign HPC or cloud partner is built for.

<200 ms

target latency for every memory read on the agent hot path

~3M

MCP calls per day at full scale: 30k clients across three verticals

1.5–3 TB

RAM for the graph layer, before hot/cold tiering

REQUIREMENT	WHY IT MATTERS
High-RAM nodes · 256 GB+ · NVMe	Memgraph keeps active tenant graphs in memory for sub-200 ms reads; RAM, not CPU, is the dominant constraint
GPU · H100 or L40S	Self-hosted inference for knowledge extraction (7–30B models) removes external API quotas
Private network · static IPs	Low-latency shard replication; MCP clients pin their endpoints
EU residency · sovereign options	Enterprise and public-sector clients need jurisdiction control over their knowledge graphs
S3-compatible object storage	100–150 TB for document corpus, OCR cache and graph snapshots

Already working with: LuxProvide / MeluXina (HPC) · NVIDIA · Amazon Web Services. Flat, predictable pricing is preferred.

PARTNERSHIP MODELS

Technology *& infrastructure*

Two partner tracks — pick one or combine them.

Technology

Agent & developer platforms

Agent, IDE and workflow platforms embed Membria memory through MCP and A2A — your users get persistent, auditable experience out of the box.

Technology

AI, data & tooling vendors

Connect your models, data pipelines or engineering tools to the shared knowledge graph: memory, verification and standards as a service inside your product.

Infrastructure

Cloud & HPC providers

Host Membria for sovereign, self-hosted and on-premise deployments — with a joint route to enterprise and public-sector buyers.

Infrastructure

GPU & hardware vendors

Co-optimise graph and inference workloads for your platform and publish reference deployments for AI-agent memory at scale.

Flexible by design: commercial terms are agreed per partner. Start with a pilot, scale what works.

MOAT

Three EMPIRIA graphs, *day one*

Cold start is the hardest problem in agent memory. We don't have one — every vertical ships with a pre-loaded PK/NK library before a single client session runs.

VERTICAL	PK — BEST PRACTICES	NK — ANTI-PATTERNS / PROHIBITIONS	SOURCE	STATUS
Engineering (EPC)	**~1.3M**	**~1.2M**	20,000+ standards — ISO · IEC · ASTM · ГОСТ · API · ASME · NFPA · DNV	LIVE · SCALING · EPC.MEMBRIA.AI
Code	**~250K Skills**	**~320K AntiPatterns**	96.5K mined repos + curated sources, CWE-mapped	LIVE · SCALING · CODE.MEMBRIA.AI
Audit / Glass	**schema live**	**schema live**	materiality & provenance graph — seeding from client engagements	DESIGN-PARTNER STAGE

Same schema, same loop, three domains. A generic memory startup starts every client at zero. We start every client already knowing what fails.

HOW WE START

From first call *to co-launch*

A short, low-risk path — you see the value before committing.

01

Align

Share your platform or infrastructure and the customers you serve; we map where memory, verification or knowledge add value.

02

Pilot

Connect with one command and run a scoped pilot on your own stack, with our engineers alongside.

03

Integrate

Embed Membria or deploy it on your infrastructure; agree deployment model, support and security review.

04

Co-launch

Go to market together: reference deployments, joint case studies and a shared pipeline.

Start in minutes: `npx @membria/cli init` — then talk to us at **hi@actiq.ai**



Let's build the layer
agents can't work without.

PARTNER WITH US

Let's Build Together



Michael Aprossine
CEO · Co-Founder
Platform architecture, enterprise delivery



Philippe Khomenok
COO · Co-Founder
Strategic vision, go-to-market and growth



Mike Keer
CTO
Memory & reasoning systems — fast, reliable, built to scale



Mayya Aprosina
UI/UX
Product design across CE, EPC and Code surfaces

PARTNER WITH US	DETAIL
Who we look for	technology: agent & developer platforms, AI, data and tooling vendors · infrastructure: cloud, HPC, GPU and hardware providers
Proof	1 bank + 1 government live enterprise · 100+ seats deployed · technology partners NVIDIA · AWS · LuxProvide
Next step	30-minute intro call, then a scoped pilot
Contact	hi@actiq.ai · membria.ai · code.membria.ai · epc.membria.ai