MEMBRIA.AI · PARTNERSHIPS

The experience layer for *AI reasoning*.

Membria records what agents decide and what happened next, so every agent improves without changing model weights or grading itself — an auditable alternative to recursive self-improvement.

We partner with technology and infrastructure providers to bring it to more teams, clouds and data centres.



Michael Aprossine

CEO, Co-Founder — Membria.AI by Actiquest Labs

PARTNER DECK

2026 · v1.0

COMPANY

Who We Are



Actiquest Labs

Applied AI Systems Company

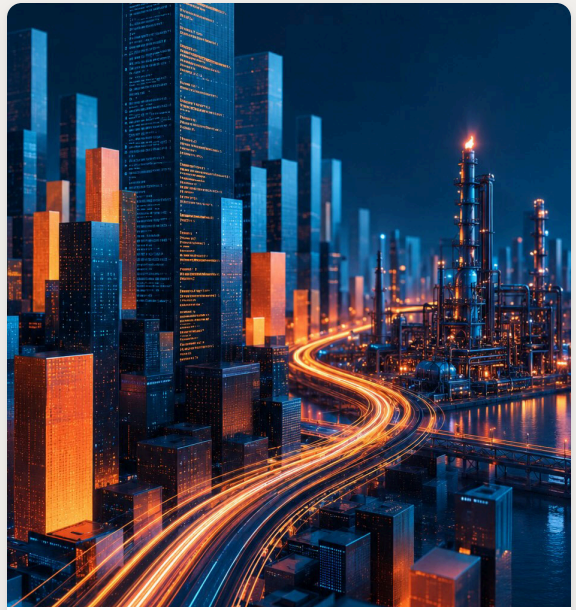
Practical automation, intelligent workflows and enterprise-grade AI — built for real results.



Membria.AI

AI Operating Environment

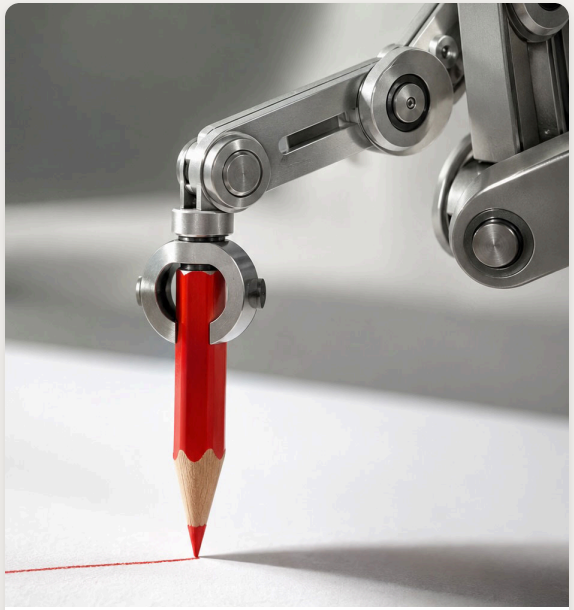
Our core product — agents, memory, knowledge and verification combined into one enterprise-ready AI platform.



Vertical Expertise

Code · Fintech · Government · Audit · EPC

We specialize in complex domains where AI needs precision, context and process understanding.



How We Work

Applied AI Engineering

We ship enterprise AI projects that stress-test and fund Membria's core — concept to production.

COMPANY PROFILE

Actiquest Labs at a Glance

Key facts, production footprint and enterprise AI capabilities.

10+ years enterprise product, AI engineering and venture-building

4 verticals — Fintech · Government · Audit · Engineering/EPC

100+ seats deployed across production AI use cases

METRIC	VALUE
Founded	February 2023 · California, USA
European Presence	Belgium entity · relocation to Luxembourg planned
Core Product	Membria.AI — memory & experience layer for AI agents
Production Verticals	SDLC · Fintech/Banking · AI Accounting/Audit · Engineering/EPC
Deployment	Cloud SaaS · self-hosted · on-premise enterprise
Clients & Partners	OES · Wowcube · Hiscox · K-Telco · Kazakhstan Government · industrial enterprises across EMEA/CIS
Technology Partners	NVIDIA · Amazon Web Services · LuxProvide
Team	4 founders + 10+ developers · 20+ years enterprise software & audit · 9 years applied AI — finance, government, industrial, deep-tech

OPPORTUNITY

Why partner *with Membria*

1. **Agents became the workforce**
   Claude Code, Codex and Cursor do real daily work in code, engineering and finance — but they ship without memory. Membria adds it.

2. **One integration, every agent**
   MCP plugs a memory plane into any agent on the market — technology partners connect once instead of building per-vendor integrations.

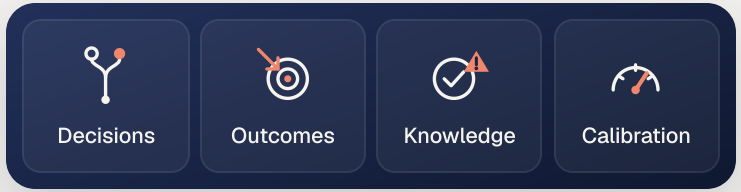
3. **Sovereign & regulated by design**
   EU AI Act and ISO 19650 require defensible decisions and audit trails. Membria runs cloud, self-hosted or on-premise — a fit for sovereign and regulated infrastructure.

4. **Experience compounds with every partner**
   Nobody owns the decision → outcome graph yet. Every deployment strengthens the shared knowledge graph that all partners build on.

PRODUCT

Membria, *the experience layer*

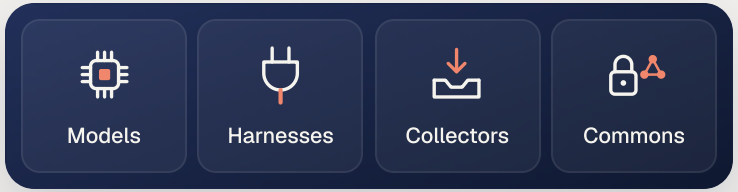
Judgment = reasoning from your model + experience from Membria. The reasoning is theirs. The experience is ours.



Experience

What is kept

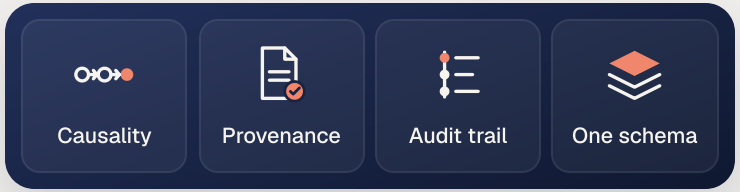
- ✓ **Decisions** — captured as made, with alternatives and confidence
- ✓ **Outcomes** — come from reality, not the model, linked to the decision
- ✓ **Knowledge** — failures become warnings, wins become practice
- ✓ **Calibration** — where the model is reliable, measured per domain



Integration

How it plugs in

- ✓ **Models** — any model, frontier or local; the experience is the same
- ✓ **Harnesses** — any harness, over one MCP endpoint, nothing to install
- ✓ **Collectors** — outcomes from git, CI, QA, audit, incident review
- ✓ **Commons** — your graph stays private, anti-patterns are shared



Foundation

Why it can be trusted

- ✓ **Causality** — decision linked to consequence, not a vector store
- ✓ **Provenance** — every fact has a source, a standard, an approver
- ✓ **Audit trail** — who decided what, when, on what evidence
- ✓ **One schema** — code, engineering, finance, emergency, music

HOW IT WORKS

One loop, *every decision*

Reasoning goes in. Experience comes out. Every real outcome feeds the next decision.

01

Recall

Surfaces similar past decisions, how they ended, and how reliable the model is in this domain.

past decisions · failures · calibration

02

Decide

The choice is captured as it is made: the statement, the alternatives, the stated confidence.

statement · alternatives · confidence

03

Outcome

The result arrives from reality, not from the model, and is linked to the decision that caused it.

revert · audit · incident · survival

04

Learn

Failure becomes a warning, success becomes practice, and confidence is recalibrated per domain.

warning · practice · recalibration

EXPERIENCE FEEDS THE NEXT DECISION

WHY MEMBRIA?

Core Capabilities

The full stack — from memory and reasoning to deployment and governance



AI Agent Orchestration

Multi-agent BBS mesh, WorkGraph DAG, A2A protocol coordination



Causal Memory

Decision / Outcome / NegativeKnowledge graph with time-decay and compounding



Verification Engine

Cross-document consistency checks, rule-based normative validation, quality gates



Graph Intelligence

Memgraph property graph, symbol-level CodeGraph, policy graph with Hebbian co-activation



LLM Integration

Claude / GPT factory, structured output, behavior chains, context injection



Product Engineering

Full-stack: FastAPI + React + Node.js, MCP protocol, Telegram, Web Dashboard



Domain AI Systems

Vertical-specific schemas, critics, standards compliance (ISA / API / ASME / IEC)



Security & Governance

PolicyEngine, RBAC, OIDC/SSO, prompt injection defense, secrets management

Bottom line: Generic AI forgets. Membria remembers, reasons, and governs.

PROJECTS

What we are *building*

Six projects on one memory core — from developer tools to regulated finance and public safety.

PROJECT	CUSTOMER / DOMAIN	WHAT WE BUILT	STATUS
Membria Code	solo developers · Claude Code, Codex, Cursor (API to B2B)	remembers every decision and what happened next, so coding agents stop repeating mistakes	SAAS
Membria CLI	engineering teams · inside the customer perimeter	one delivery pipeline for a team's agents, from spec to release, with gates, policy and audit	ON-PREMISE
AI-EPC	EPC contractors & design institutes · Worldwide	turns a tender folder into a checked, issued engineering package across 17 disciplines	SAAS
Broker Platform	licensed virtual-asset broker · Central Asia	data bus and decision memory that connect execution and compliance	ON-PREMISE
Bank automation	developer-owned bank · CIS	regulatory report specs, mortgage checks and escrow rules, drafted and validated by agents	ON-PREMISE
C2 Center AI	Government Emergency Response · Central Asia	early warning and road-closure control, with every decision approved by a person	ON-PREMISE

Two delivery models: SaaS to start in minutes, On-Premise inside the customer's perimeter. Most client names are withheld.

USE CASE: SAAS

Membria Code *learns from reverts*

Solo developers · Claude Code, Codex, Cursor (API to B2B) · SaaS

A hosted experience layer for developers who ship with Claude Code, Codex or Cursor. One MCP endpoint, a private graph per account, nothing to run.

1. The situation

A developer asks Claude Code to add caching to an API client. The model reaches for an in-memory LRU — the same choice that broke under multi-worker deploys and was reverted two weeks ago. That lesson lives in an old transcript nobody will read.

2. How it works

Before choosing, the model asks Membria for decision context. It gets back the reverted decision and its negative knowledge — per-process cache, inconsistent across workers, recorded with the decision — plus the commons signal: how often this pattern is removed in other projects.

3. The result

The model picks a shared cache and records the decision with its alternatives. Two weeks later the change is still in place, and the outcome attaches itself from git — nobody writes a post-mortem. The model is unchanged and never grades its own work: the outcome comes from reality, and the next decision starts from what actually happened.

96.5K

production repos mined

570K

skills + anti-patterns in the commons

1

command to connect

```
npx @membria/cli init
```

Works with: Claude Code · Codex · Cursor · any MCP client · code.membria.ai

CASE STUDY · 02

Membria CLI *agents as one team*

Engineering teams · inside the customer perimeter · On-Premise

An SDLC harness that runs a team's agents as one delivery pipeline, from spec to release, with gates, policy and audit, and starts from the experience accumulated in Membria Code.

1. The situation

A team of six has to add rate limiting to the public API. Instead of six separate agent sessions, the task enters one pipeline: Researcher, Planner, Developer, Reviewer and Gate.

2. How it works

The Researcher surfaces a per-process counter the team shipped and reverted last year, and Membria Code adds what the team has never seen: in most projects that adopted this pattern, it was later removed. The Planner builds the plan against CodeGraph and adds a load test. An engineer's own Claude Code takes the Developer seat over MCP.

3. The result

The Gate blocks the first attempt for missing load-test evidence and passes the second. Two weeks later the change has survived, and the outcome becomes the team's own experience for the next task. Every step has an actor, a model, a policy decision and an audit record.

570K

skills and anti-patterns available from day one

5

agent roles, one blocking gate

0

bytes of the team graph leave the perimeter

Stack & partners: MCP · OIDC / SSO · OPA · any LLM provider · Membria Code commons (read-only)

CASE STUDY · 03

AI-EPC *from tender to issue*

EPC contractors & design institutes · Worldwide · SaaS

1. The situation

An EPC contractor uploads a tender folder: specifications, drawings, vendor catalogues, contract terms. This is where requirements get missed, and a late clause turns into rework — 28% of budget overruns.

2. How it works

AI-EPC reads every file, scanned pages included, into a requirements register where each line cites its source clause. It builds the design level by level: models propose, rules, codes and catalogues accept or reject, and unknown values become questions to the customer.

3. The result

Each discipline lead approves documents at a gate, and a change in one place reaches every dependent sheet. The contractor issues a consistent package with an evidence trail behind each document. Engineers sign; the system never stamps.

17
engineering disciplines

20,000+
standards in the graph

5
approval gates, G0 to G4

Stack & partners: DXF · DEXPI · IFC · ISO / IEC / ASME / API / GOST · epc.membria.ai

CASE STUDY · 04

Broker Platform *one data bus*

Licensed virtual-asset broker · Central Asia · client anonymised · On-Premise

1. The situation

A licensed broker runs execution and compliance as two parallel tracks built by two vendors. An operation must not execute until compliance has decided on it, and personal data must stay out of shared systems.

2. How it works

Membria owns the data bus and decision memory between the tracks: a Kafka backbone with 38 schema-validated event types and a graph of 42 object types. A client's status and every operation decision travel to the other track over the bus.

3. The result

A compliance officer asks in plain language why an operation was blocked and gets the full chain — screening result, limit breaches, approval steps — with the source events. Both vendors integrate against one documented contract.

38

event types on one bus

376

automated tests passing

~6 wks

from design to first full cycle

Stack & partners: Kafka · Memgraph · Qdrant · FastAPI · Python

CASE STUDY · 05

Bank *reporting & mortgage*

Developer-owned bank · CIS · regulatory reporting, mortgage and escrow · On-Premise

1. The situation

Methodologists spend 2–5 business days on each central-bank report specification. Mortgage review is a 40+ field manual checklist, and escrow specs are copied from the previous project.

2. How it works

A reporting assistant writes the specification from a free-form task, using graphs of the bank's object model and the regulator's forms. Skills keep the mappings that worked; rejected mappings are never proposed again. A mortgage assistant checks each application against project and escrow data.

3. The result

A specification is ready in 1–4 hours instead of 2–5 days, and when the regulator changes a form, every affected specification is listed at once. The bank picks its own model — Claude, GigaChat or on-prem Qwen — and data stays inside its security boundary.

1–4 h

per specification, down from 2–5 days

3

rollout phases, about 20 weeks

4

domain graphs: model, regulatory, compliance, escrow

Stack & partners: Graph memory · Claude / GigaChat / on-prem Qwen · web, Telegram, CLI

CASE STUDY · 06

C2 Center: *Palantir-class, ministry-owned*

Ministry of Emergency Situations, Kazakhstan · National command system · In delivery

1. **The challenge**

   Emergency data sat in separate systems, and decisions depended on manual cross-checks. The ministry needed AI that would speed up command decisions without data leaving the country or decisions leaving human hands.

2. **What we built**

   A command platform where five AI agents assess risk, draft response plans, build routes, prepare alerts and analyse outcomes. Every recommendation comes with its sources and a confidence score. A quality gate blocks plans that repeat past mistakes, and the system learns from every incident.

3. **Why it stands out**

   It runs entirely inside the ministry's network on its own AI models, and the ministry keeps the source code. Humans approve every action, in line with Kazakhstan's new AI law.

3 years
national program

Minutes
from alert to response plan

0
data leaves the country

100%
ministry-owned source code

Key tech: knowledge graph · on-premise LLMs · explainable AI · audit trail

TEAM

The people *behind it*

Four founders, two decades in IT, IoT, fintech and applied AI — and a team of 10+ developers.



Michael Aprossine

CEO · CO-FOUNDER

Serial founder with 20+ years of product leadership across enterprise IT, IoT, mobility and blockchain. Started in IT consulting at a Microsoft Gold Partner, co-founded Apla/eGaaS and sold the mobility startup DrivePoint to Hyundai. Owns Membria's product vision and system architecture.

Focus: product vision, system architecture, AI platform
Previously: Apla/eGaaS · DrivePoint (now part of Hyundai Mobility) · Microsoft Business Solutions/Gold Partner

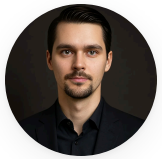


Philippe Khomenok

COO · CO-FOUNDER

Commercial builder who turns new markets into revenue. Scaled online sales of cable and rolled-metal products to $10M a year from scratch, co-founded the Grown Capital fund ($5M raised) and, as CBDO at Advcash, launched new products and entered new markets. Leads Membria's go-to-market and partnerships.

Focus: go-to-market, partnerships, growth
Previously: Advcash (CBDO) · Grown Capital (co-founder)



Mike Keer

CTO

AI architect with deep expertise in language models and computer vision. Leads Membria's engineering and designs its memory and reasoning systems. Has worked on joint projects with Luxinnovation, the European Innovation Council and others.

Focus: memory & reasoning systems, LLMs, computer vision
Previously: joint projects with Luxinnovation and the EIC



Mayya Vergouwen

UI/UX

Product designer with 6+ years leading cross-functional teams through the full product cycle — research, journey mapping, design systems and front-end build in React and Next.js. Founder of the SiteSynth web studio and the designer behind ScoreSynth and Membria.

Focus: UI/UX, design systems, front-end
Previously: SiteSynth (founder) · Idalko (digital specialist partner)

INFRASTRUCTURE PARTNERS

What we need *from infrastructure*

Agent memory is RAM-bound and latency-critical — the profile a sovereign HPC or cloud partner is built for.

<200 ms

target latency for every memory read on the agent hot path

~3M

MCP calls per day at full scale: 30k clients across three verticals

1.5–3 TB

RAM for the graph layer, before hot/cold tiering

REQUIREMENT	WHY IT MATTERS
High-RAM nodes · 256 GB+ · NVMe	Memgraph keeps active tenant graphs in memory for sub-200 ms reads; RAM, not CPU, is the dominant constraint
GPU · H100 or L40S	Self-hosted inference for knowledge extraction (7–30B models) removes external API quotas
Private network · static IPs	Low-latency shard replication; MCP clients pin their endpoints
EU residency · sovereign options	Enterprise and public-sector clients need jurisdiction control over their knowledge graphs
S3-compatible object storage	100–150 TB for document corpus, OCR cache and graph snapshots

Already working with: LuxProvide / MeluXina (HPC) · NVIDIA · Amazon Web Services. Flat, predictable pricing is preferred.

HOW WE START

From first call *to co-launch*

A short, low-risk path: you see measured results before you commit.

01

Align

Week 1

Share your platform and the users you serve. We pick the one workflow where repeated mistakes cost the most and agree on how to measure it.

YOU GET

a pilot scope and a baseline metric.

02

Pilot

4–6 weeks

Connect over MCP or API and run Membria on that workflow with real users, our engineers alongside. Outcomes are collected automatically from your own systems.

YOU GET

before-and-after numbers on repeated mistakes, reverts or rework, plus a go/no-go decision based on them.

03

Integrate

4–8 weeks

Embed Membria in your product, hosted or on your infrastructure. Agree the deployment model, data boundaries, support and security review.

YOU GET

a production integration with a private graph for each of your customers.

04

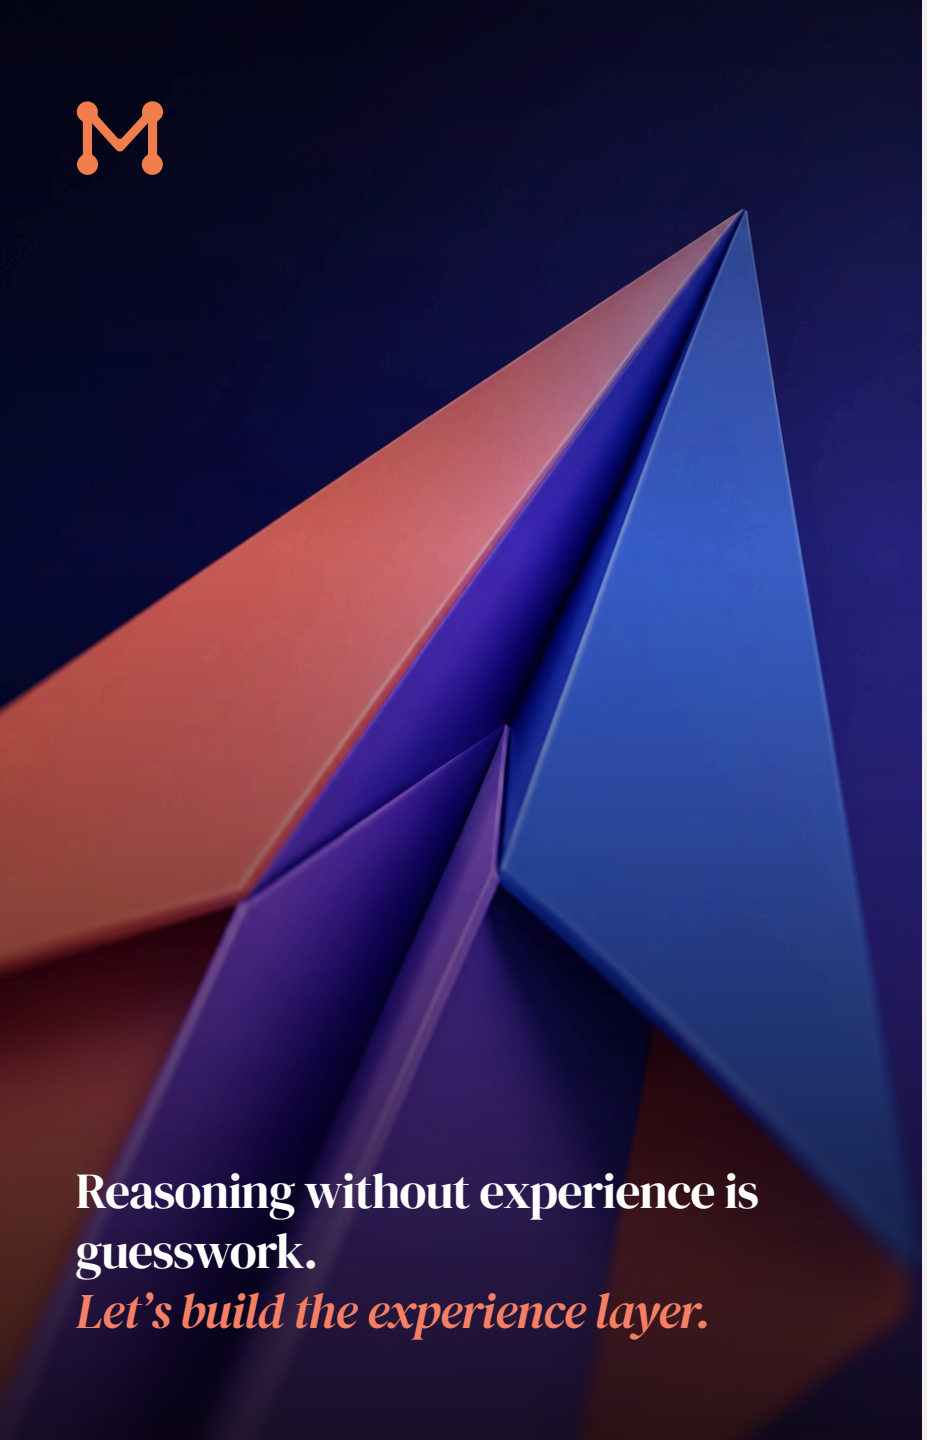
Co-launch

Offer it to your users together: reference deployments, a joint case study built on the pilot's numbers, and a shared pipeline.

YOU GET

a published case and joint go-to-market.

Partnerships: hi@actiq.ai



Reasoning without experience is guesswork.
Let's build the experience layer.

PARTNER WITH US

Let's build together

What we look for	Corporate pilots: enterprises that want AI agents to remember, verify and audit their decisions — banks, EPC and engineering firms, government agencies · Partners who bring us to them: AI coding and agent platforms, data and tooling vendors, cloud, HPC and GPU providers
Track record	Deployed with a bank and a government agency · 100+ seats in use · 10K SaaS users · Working with NVIDIA, AWS and LuxProvide
Next step	A 30-minute intro call, then a scoped pilot with one measurable outcome
Contact	hi@actiq.ai · membria.ai · code.membria.ai · epc.membria.ai