MEMBRIA.AI · PARTNERSHIPS

# The experience layer for *AI reasoning*.

Membria records what agents decide and what happened next, so every agent improves without changing model weights or grading itself — an auditable alternative to recursive self-improvement.

We partner with technology and infrastructure providers to bring it to more teams, clouds and data centres.



**Michael Aprossine**

CEO, Co-Founder — Membria.AI by Actiquest Labs

PARTNER DECK

2026 · v1.0

COMPANY

# Who We Are



## Actiquest Labs

**Applied AI Systems Company**

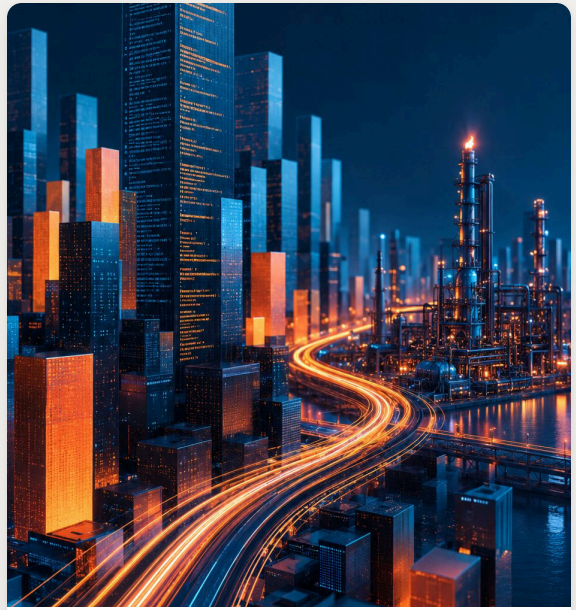
Practical automation, intelligent workflows and enterprise-grade AI — built for real results.



## Membria.AI

**AI Operating Environment**

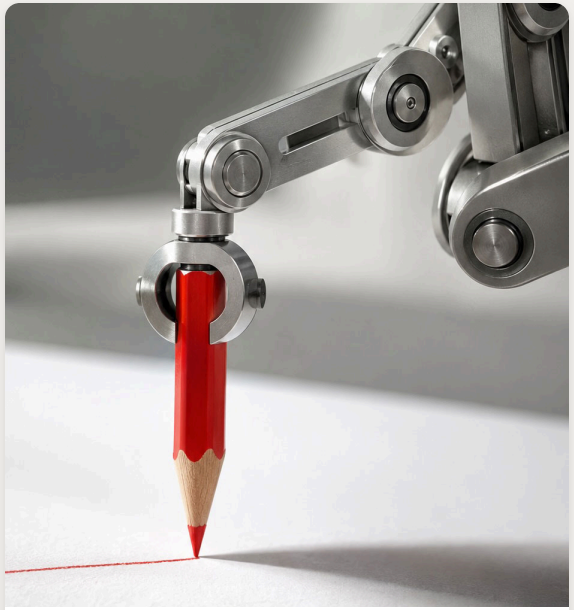
Our core product — agents, memory, knowledge and verification combined into one enterprise-ready AI platform.



## Vertical Expertise

**Code · Fintech · Government · Audit · EPC**

We specialize in complex domains where AI needs precision, context and process understanding.



## How We Work

**Applied AI Engineering**

We ship enterprise AI projects that stress-test and fund Membria's core — concept to production.

COMPANY PROFILE

# Actiquest Labs at a Glance

Key facts, production footprint and enterprise AI capabilities.

**10+** years enterprise product, AI engineering and venture-building

**4** verticals — Fintech · Government · Audit · Engineering/EPC

**100+** seats deployed across production AI use cases

| METRIC | VALUE |
|---|---|
| **Founded** | February 2023 · California, USA |
| **European Presence** | Belgium entity · relocation to Luxembourg planned |
| **Core Product** | Membria.AI — memory & experience layer for AI agents |
| **Production Verticals** | SDLC · Fintech/Banking · AI Accounting/Audit · Engineering/EPC |
| **Deployment** | Cloud SaaS · self-hosted · on-premise enterprise |
| **Clients & Partners** | OES · Wowcube · Hiscox · K-Telco · Kazakhstan Government · industrial enterprises across EMEA/CIS |
| **Technology Partners** | NVIDIA · Amazon Web Services · LuxProvide |
| **Team** | core team of 4 · 20+ years enterprise software & audit · 9 years applied AI — finance, government, industrial, deep-tech |

OPPORTUNITY

# Why partner *with Membria*

1. **Agents became the workforce**
   Claude Code, Codex and Cursor do real daily work in code, engineering and finance — but they ship without memory. Membria adds it.

2. **One integration, every agent**
   MCP plugs a memory plane into any agent on the market — technology partners connect once instead of building per-vendor integrations.

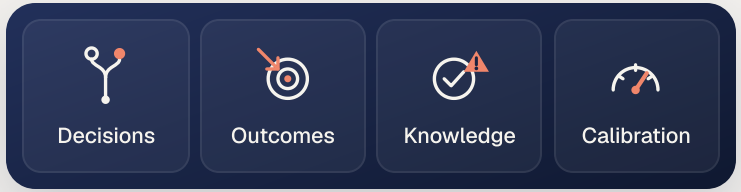
3. **Sovereign & regulated by design**
   EU AI Act and ISO 19650 require defensible decisions and audit trails. Membria runs cloud, self-hosted or on-premise — a fit for sovereign and regulated infrastructure.

4. **Experience compounds with every partner**
   Nobody owns the decision → outcome graph yet. Every deployment strengthens the shared knowledge graph that all partners build on.

PRODUCT

# Membria, *the experience layer*

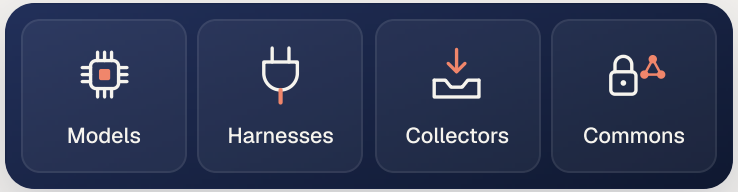
Judgment = reasoning from your model + experience from Membria. The reasoning is theirs. The experience is ours.




## Experience

What is kept

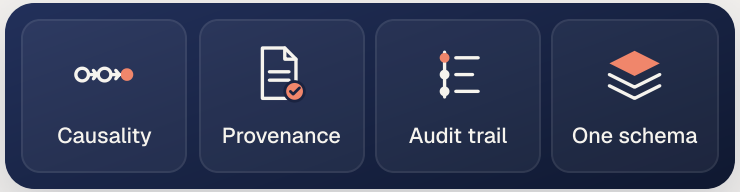
- ✓ **Decisions** — captured as made, with alternatives and confidence
- ✓ **Outcomes** — come from reality, not the model, linked to the decision
- ✓ **Knowledge** — failures become warnings, wins become practice
- ✓ **Calibration** — where the model is reliable, measured per domain




## Integration

How it plugs in

- ✓ **Models** — any model, frontier or local; the experience is the same
- ✓ **Harnesses** — any harness, over one MCP endpoint, nothing to install
- ✓ **Collectors** — outcomes from git, CI, QA, audit, incident review
- ✓ **Commons** — your graph stays private, anti-patterns are shared




## Foundation

Why it can be trusted

- ✓ **Causality** — decision linked to consequence, not a vector store
- ✓ **Provenance** — every fact has a source, a standard, an approver
- ✓ **Audit trail** — who decided what, when, on what evidence
- ✓ **One schema** — code, engineering, finance, emergency, music

HOW IT WORKS

# One loop, *every decision*

Reasoning goes in. Experience comes out. Every real outcome feeds the next decision.

*01*

### Recall

Surfaces similar past decisions, how they ended, and how reliable the model is in this domain.

past decisions · failures · calibration

*02*

### Decide

The choice is captured as it is made: the statement, the alternatives, the stated confidence.

statement · alternatives · confidence

*03*

### Outcome

The result arrives from reality, not from the model, and is linked to the decision that caused it.

revert · audit · incident · survival

*04*

### Learn

Failure becomes a warning, success becomes practice, and confidence is recalibrated per domain.

warning · practice · recalibration

EXPERIENCE FEEDS THE NEXT DECISION

WHY MEMBRIA?

# Core Capabilities

The full stack — from memory and reasoning to deployment and governance



### AI Agent Orchestration

Multi-agent BBS mesh, WorkGraph DAG, A2A protocol coordination



### Causal Memory

Decision / Outcome / NegativeKnowledge graph with time-decay and compounding



### Verification Engine

Cross-document consistency checks, rule-based normative validation, quality gates



### Graph Intelligence

Memgraph property graph, symbol-level CodeGraph, policy graph with Hebbian co-activation



### LLM Integration

Claude / GPT factory, structured output, behavior chains, context injection



### Product Engineering

Full-stack: FastAPI + React + Node.js, MCP protocol, Telegram, Web Dashboard



### Domain AI Systems

Vertical-specific schemas, critics, standards compliance (ISA / API / ASME / IEC)



### Security & Governance

PolicyEngine, RBAC, OIDC/SSO, prompt injection defense, secrets management

**Bottom line:** Generic AI forgets. Membria remembers, reasons, and governs.

PROJECTS

# What we are *building*

Six projects on one memory core — from developer tools to regulated finance and public safety.

| PROJECT | CUSTOMER / DOMAIN | WHAT WE BUILT | STATUS |
|---|---|---|---|
| **Membria CE** | solo developers · Claude Code, Codex, Cursor (API to B2B) | remembers every decision and what happened next, so coding agents stop repeating mistakes | SAAS |
| **Membria CLI** | engineering teams · inside the customer perimeter | one delivery pipeline for a team's agents, from spec to release, with gates, policy and audit | ON-PREMISE |
| **AI-EPC** | EPC contractors & design institutes · Worldwide | turns a tender folder into a checked, issued engineering package across 17 disciplines | SAAS |
| **Broker Platform** | licensed virtual-asset broker · Central Asia | data bus and decision memory that connect execution and compliance | ON-PREMISE |
| **Bank automation** | developer-owned bank · CIS | regulatory report specs, mortgage checks and escrow rules, drafted and validated by agents | ON-PREMISE |
| **C2 Center AI** | Government Emergency Response · Central Asia | early warning and road-closure control, with every decision approved by a person | ON-PREMISE |

**Two delivery models:** SaaS to start in minutes, On-Premise inside the customer's perimeter. Client names are withheld.

USE CASE · MEMBRIA CODE SAAS

# Membria CE *learns from reverts*

**Solo developers · Claude Code, Codex, Cursor (API to B2B) · SaaS**

A hosted experience layer for developers who ship with Claude Code, Codex or Cursor. One MCP endpoint, a private graph per account, nothing to run.

## 1. The situation

A developer asks Claude Code to add caching to an API client. The model reaches for an in-memory LRU — the same choice that broke under multi-worker deploys and was reverted two weeks ago. That lesson lives in an old transcript nobody will read.

## 2. How it works

Before choosing, the model asks Membria for decision context. It gets back the reverted decision and its negative knowledge — per-process cache, inconsistent across workers, recorded with the decision — plus the commons signal: how often this pattern is removed in other projects.

## 3. The result

The model picks a shared cache and records the decision with its alternatives. Two weeks later the change is still in place, and the outcome attaches itself from git — nobody writes a post-mortem. The model is unchanged and never grades its own work: the outcome comes from reality, and the next decision starts from what actually happened.

**96.5K**

production repos mined

**570K**

skills + anti-patterns in the commons

**1**

command to connect

```
npx @membria/cli init
```

**Works with:** Claude Code · Codex · Cursor · any MCP client · code.membria.ai

CASE STUDY · 02

# Membria CLI *agents as one team*

Engineering teams · inside the customer perimeter · On-Premise

**1. The situation**

A team of six has to add rate limiting to the public API. Instead of six separate agent sessions, the task enters one pipeline: Researcher, Planner, Developer, Reviewer and Gate.

**2. How it works**

The Researcher surfaces a per-process counter the team shipped and reverted last year. The Planner builds the plan against CodeGraph and adds a load test. An engineer's own Claude Code takes the Developer seat over MCP.

**3. The result**

The Gate blocks the first attempt for missing load-test evidence and passes the second. Every step has an actor, a model, a policy decision and an audit record, and the team graph never leaves the perimeter.

**8**

pipeline phases, spec to release

**5**

agent roles, one blocking gate

**100%**

of the team graph stays in the perimeter

**Stack & partners:** MCP · OIDC / SSO · OPA · any LLM provider

CASE STUDY · 03

# AI-EPC *from tender to issue*

EPC contractors & design institutes · Worldwide · SaaS

**1. The situation**

An EPC contractor uploads a tender folder: specifications, drawings, vendor catalogues, contract terms. This is where requirements get missed, and a late clause turns into rework — 28% of budget overruns.

**2. How it works**

AI-EPC reads every file, scanned pages included, into a requirements register where each line cites its source clause. It builds the design level by level: models propose, rules, codes and catalogues accept or reject, and unknown values become questions to the customer.

**3. The result**

Each discipline lead approves documents at a gate, and a change in one place reaches every dependent sheet. The contractor issues a consistent package with an evidence trail behind each document. Engineers sign; the system never stamps.

**17**

engineering disciplines

**20,000+**

standards in the graph

**5**

approval gates, G0 to G4

**Stack & partners:** DXF · DEXPI · IFC · ISO / IEC / ASME / API / GOST · epc.membria.ai

CASE STUDY · 04

# Broker Platform *one data bus*

Licensed virtual-asset broker · Central Asia · client anonymised · On-Premise

1. **The situation**
   A licensed broker runs execution and compliance as two parallel tracks built by two vendors. An operation must not execute until compliance has decided on it, and personal data must stay out of shared systems.

2. **How it works**
   Membria owns the data bus and decision memory between the tracks: a Kafka backbone with 38 schema-validated event types and a graph of 42 object types. A client's status and every operation decision travel to the other track over the bus.

3. **The result**
   A compliance officer asks in plain language why an operation was blocked and gets the full chain — screening result, limit breaches, approval steps — with the source events. Both vendors integrate against one documented contract.

**38**
event types on one bus

**376**
automated tests passing

**~6 wks**
from design to first full cycle

**Stack & partners:** Kafka · Memgraph · Qdrant · FastAPI · Python

CASE STUDY · 05

# Bank *reporting & mortgage*

Developer-owned bank · CIS · regulatory reporting, mortgage and escrow · On-Premise

## 1. The situation

Methodologists spend 2–5 business days on each central-bank report specification. Mortgage review is a 40+ field manual checklist, and escrow specs are copied from the previous project.

## 2. How it works

A reporting assistant writes the specification from a free-form task, using graphs of the bank's object model and the regulator's forms. Skills keep the mappings that worked; rejected mappings are never proposed again. A mortgage assistant checks each application against project and escrow data.

## 3. The result

A specification is ready in 1–4 hours instead of 2–5 days, and when the regulator changes a form, every affected specification is listed at once. The bank picks its own model — Claude, GigaChat or on-prem Qwen — and data stays inside its security boundary.

**1–4 h**

per specification, down from 2–5 days

**3**

rollout phases, about 20 weeks

**4**

domain graphs: model, regulatory, compliance, escrow

**Stack & partners:** Graph memory · Claude / GigaChat / on-prem Qwen · web, Telegram, CLI

CASE STUDY · 06

# C2 Center *decisions in minutes*

Government Emergency Response · Central Asia · On-Premise

1. **The alert**
   Drone LiDAR shows slope movement near a bridge well above the seasonal limit, while river levels rise. The duty officer needs one decision-ready picture, not five separate systems.

2. **How it works**
   The Duty agent matches the finding with the object's history in the knowledge graph and raises a risk alert with a confidence score and cited sources. The Planner offers three response options, Routes builds detours, and quality control blocks any option that repeats a recorded unsafe decision.

3. **The result**
   The chief of staff approves, and only then do alerts go out in two languages by SMS, Cell Broadcast and sirens. On closed roads, edge-AI cameras send plate events instead of video, with no vehicle-owner registry; the full trail stays on-premise.

**5**
AI agents, humans decide

**100%**
on-premise processing

**0**
personal-data registries

**Stack & partners:** LiDAR · IoT sensors · MQTT with mutual TLS · edge-AI cameras

TEAM

# The people *behind it*

Founders and operators with two decades across enterprise IT, IoT, fintech and applied AI.



## Michael Aprossine

**CEO · CO-FOUNDER**

Serial founder with 20+ years of product leadership across enterprise IT, IoT, mobility and blockchain. Started in IT consulting at a Microsoft Gold Partner, co-founded Apla/eGaaS and sold the mobility startup DrivePoint to Hyundai. Owns Membria's product vision and system architecture.

**Focus:** product vision, system architecture, AI platform
**Previously:** Apla/eGaaS · DrivePoint (now Hyundai Mobility) · Slotlogic IoT/SaaS · Microsoft Gold Partner

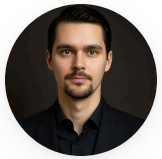


## Philippe Khomenok

**COO · CO-FOUNDER**

Commercial builder who turns new markets into revenue. Scaled online sales of cable and rolled-metal products to $10M a year from scratch, co-founded the Grown Capital fund ($5M raised) and, as CBDO at Advcash, launched new products and entered new markets. Leads Membria's go-to-market and partnerships.

**Focus:** go-to-market, partnerships, growth
**Previously:** Advcash (CBDO) · Grown Capital (co-founder)



## Mike Keer

**CTO**

AI architect with deep expertise in language models and computer vision. Leads Membria's engineering and designs its memory and reasoning systems. Has worked on joint projects with Luxinnovation, the European Innovation Council and others.

**Focus:** memory & reasoning systems, LLMs, computer vision
**Previously:** joint projects with Luxinnovation and the EIC



## Mayya Aprosina

**UI/UX**

Product designer with 6+ years leading cross-functional teams through the full product cycle — research, journey mapping, design systems and front-end build in React and Next.js. Founder of the SiteSynth web studio and the designer behind ScoreSynth and Membria.

**Focus:** UI/UX, design systems, front-end
**Previously:** SiteSynth (founder) · Idalko (digital specialist partner)

INFRASTRUCTURE PARTNERS

# What we need *from infrastructure*

Agent memory is RAM-bound and latency-critical — the profile a sovereign HPC or cloud partner is built for.

**<200 ms**

target latency for every memory read on the agent hot path

**~3M**

MCP calls per day at full scale: 30k clients across three verticals

**1.5–3 TB**

RAM for the graph layer, before hot/cold tiering

| REQUIREMENT | WHY IT MATTERS |
|---|---|
| **High-RAM nodes · 256 GB+ · NVMe** | Memgraph keeps active tenant graphs in memory for sub-200 ms reads; RAM, not CPU, is the dominant constraint |
| **GPU · H100 or L40S** | Self-hosted inference for knowledge extraction (7–30B models) removes external API quotas |
| **Private network · static IPs** | Low-latency shard replication; MCP clients pin their endpoints |
| **EU residency · sovereign options** | Enterprise and public-sector clients need jurisdiction control over their knowledge graphs |
| **S3-compatible object storage** | 100–150 TB for document corpus, OCR cache and graph snapshots |

**Already working with:** LuxProvide / MeluXina (HPC) · NVIDIA · Amazon Web Services. Flat, predictable pricing is preferred.

PARTNERSHIP MODELS

# Technology *& infrastructure*

Two partner tracks — pick one or combine them.

*Technology*

### Agent & developer platforms

Agent, IDE and workflow platforms embed Membria memory through MCP and A2A — your users get persistent, auditable experience out of the box.

*Technology*

### AI, data & tooling vendors

Connect your models, data pipelines or engineering tools to the shared knowledge graph: memory, verification and standards as a service inside your product.

*Infrastructure*

### Cloud & HPC providers

Host Membria for sovereign, self-hosted and on-premise deployments — with a joint route to enterprise and public-sector buyers.

*Infrastructure*

### GPU & hardware vendors

Co-optimise graph and inference workloads for your platform and publish reference deployments for AI-agent memory at scale.

**Flexible by design:** commercial terms are agreed per partner. Start with a pilot, scale what works.

MOAT

# Three EMPIRIA graphs, *day one*

Cold start is the hardest problem in agent memory. We don't have one — every vertical ships with a pre-loaded PK/NK library before a single client session runs.

| VERTICAL | PK — BEST PRACTICES | NK — ANTI-PATTERNS / PROHIBITIONS | SOURCE | STATUS |
|---|---|---|---|---|
| **Engineering (EPC)** | **~1.3M** | **~1.2M** | 20,000+ standards — ISO · IEC · ASTM · ГОСТ · API · ASME · NFPA · DNV | LIVE · SCALING · EPC.MEMBRIA.AI |
| **Code** | **~250K Skills** | **~320K AntiPatterns** | 96.5K mined repos + curated sources, CWE-mapped | LIVE · SCALING · CODE.MEMBRIA.AI |
| **Audit / Glass** | **schema live** | **schema live** | materiality & provenance graph — seeding from client engagements | DESIGN-PARTNER STAGE |

**Same schema, same loop, three domains.** A generic memory startup starts every client at zero. We start every client already knowing what fails.

HOW WE START

# From first call *to co-launch*

A short, low-risk path — you see the value before committing.

*01*

### Align

Share your platform or infrastructure and the customers you serve; we map where memory, verification or knowledge add value.

*02*

### Pilot

Connect with one command and run a scoped pilot on your own stack, with our engineers alongside.

*03*

### Integrate

Embed Membria or deploy it on your infrastructure; agree deployment model, support and security review.

*04*

### Co-launch

Go to market together: reference deployments, joint case studies and a shared pipeline.

**Start in minutes:** `npx @membria/cli init` — then talk to us at **hi@actiq.ai**



Let's build the layer
*agents can't work without.*

PARTNER WITH US

# Let's Build Together



**Michael Aprossine**
CEO · Co-Founder
Platform architecture, enterprise delivery



**Philippe Khomenok**
COO · Co-Founder
Strategic vision, go-to-market and growth



**Mike Keer**
CTO
Memory & reasoning systems — fast, reliable, built to scale



**Mayya Aprosina**
UI/UX
Product design across CE, EPC and Code surfaces

| PARTNER WITH US | DETAIL |
|---|---|
| **Who we look for** | technology: agent & developer platforms, AI, data and tooling vendors · infrastructure: cloud, HPC, GPU and hardware providers |
| **Proof** | 1 bank + 1 government live enterprise · 100+ seats deployed · technology partners NVIDIA · AWS · LuxProvide |
| **Next step** | 30-minute intro call, then a scoped pilot |
| **Contact** | hi@actiq.ai · membria.ai · code.membria.ai · epc.membria.ai |